

# Optimization on deep learning

---


MATH818 RESEARCH IN APPLIED MATH

2019020356 JAEHEUN JUNG



# DNN Architecture

---



- Universal function approximator
- Combination of weighted sum and activation
- Parametrized with matrices



# Loss function

---

- Measurement for the difference between model and real data

MSE loss (Mean-Squared Error)

$$L(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i^{true} - y_i^{pred})^2$$



Cross-entropy loss

$$L(\theta) = \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^k -y_{ij}^{true} \log(y_{ij}^{pred})$$



# Gradient descent

- First-order iterative optimization




# Optimizers

- Momentum

$$v_{t+1} = \gamma v_t - \eta \nabla L(\theta_t)$$

$$\theta_{t+1} = \theta_t + v_{t+1}$$



# optimizers

---

- Adaptive learning rate

- Adagrad

$$G_{\theta}^t = \sum_{i=0}^t (\nabla L(\theta_i))^2$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{G_{\theta}^t + \epsilon}} \nabla L(\theta_t)$$

- RMSprop

$$G_{\theta}^t = \gamma G_{\theta}^{t-1} + (1 - \gamma)(\nabla L(\theta_t))^2$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{G_{\theta}^t + \epsilon}} \nabla L(\theta^t)$$



# optimizers

---

- Adam

$$\hat{m}_\theta = \frac{m_\theta^{t+1}}{1 - \beta_1^{t+1}} \text{ where } m_\theta^{t+1} = \beta_1 m_\theta^t + (1 - \beta_1) \nabla L(\theta^t)$$

$$\hat{G}_\theta = \frac{G_\theta^{t+1}}{1 - \beta_2^{t+1}} \text{ where } G_\theta^{t+1} = \beta_2 G_\theta^t + (1 - \beta_2) (\nabla L(\theta^t))^2$$

$$\theta^{t+1} = \theta^t - \eta \frac{\hat{m}_\theta}{\sqrt{\hat{G}_\theta + \epsilon}}$$



# Optimizers

---

- RAdam

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2)(\nabla L(\theta^t))^2$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \text{ where } m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla L(\theta_t)$$

$$\rho_t = \rho_\infty - 2t \frac{\beta_2^t}{1 - \beta_2^t} \text{ where } \rho_\infty = \frac{2}{1 - \beta_2} - 1$$

if  $\rho_t > 4$  then

$$\ell_t = \sqrt{\frac{1 - \beta_2^t}{v_t}}$$

$$r_t = \sqrt{\frac{(\rho_t - 4)(\rho_t - 2)\rho_\infty}{(\rho_\infty - 4)(\rho_\infty - 2)\rho_t}}$$

$\theta_t = \theta_{t-1} - \alpha_t r_t \ell_t \hat{m}_t$  where  $\alpha_t$  is step size

else  $\theta_t = \theta_{t-1} - \alpha_t \hat{m}_t$





# Batch size


# of datapoints for single update

Stochastic Gradient descent

- Single datapoint

Batch Gradient descent

- All datapoints



# Data parallelism

- Multiple workers compute gradient
- Parameter server collects gradients and update the parameter
- Update Rules:
  - SSGD and ASGD

---

**Algorithm 2:** worker  $m$ 

---

**Input** dataset  $\mathcal{X}$ , minibatch size  $\mathcal{B}$

**for**  $t = 0, 1, \dots$  **do**

    Wait to read  $\theta^{(t)}$  from parameter server;

$G_m^{(t)} := 0$ ;

**for**  $i = 1, \dots, \mathcal{B}$  **do**

        Sample data  $x_{k,i}$  from  $\mathcal{X}$ ;

$G_m^{(t)} \leftarrow G_m^{(t)} + \frac{1}{\mathcal{B}} \nabla L(x_{k,i}, \theta^{(t)})$

**end**

    Send  $G_m^{(t)}$  to parameter server

**end**



# Data parallelism

---

## SYNCHRONOUS UPDATE

---

**Algorithm 3:** SSGD parameter server

---

**Input** learning rate  $\mu_t$  and number of workers  $M$ ;  
initialize  $t \leftarrow 0$ ;  
initialize model  $\theta^{(0)}$  ;  
**repeat**  
    Send  $\theta^{(t)}$  to each workers;  
    Wait for  $G_1^{(t)}, \dots, G_M^{(t)}$  from each workers;  
     $\theta^{(t+1)} \leftarrow \theta^{(t)} - \frac{\mu_t}{N} \sum_{l=1}^M G_l^{(t)}$ ;  
     $t \leftarrow t + 1$   
**until** *converges*;

---

## ASYNCHRONOUS UPDATE

---

**Algorithm 5:** ASGD parameter server

---


**Input** learning rate  $\mu_t$  and number of workers  $M$ ;  
initialize  $t \leftarrow 0$ ;  
initialize model  $\theta^{(0)}$  ;  
**repeat**  
    Wait for  $g_m$  from any worker;  
     $\theta^{(t+1)} \leftarrow \theta^{(t)} - \mu_t g_m$ ;  
     $t \leftarrow t + 1$   
**until** *converges*;

---

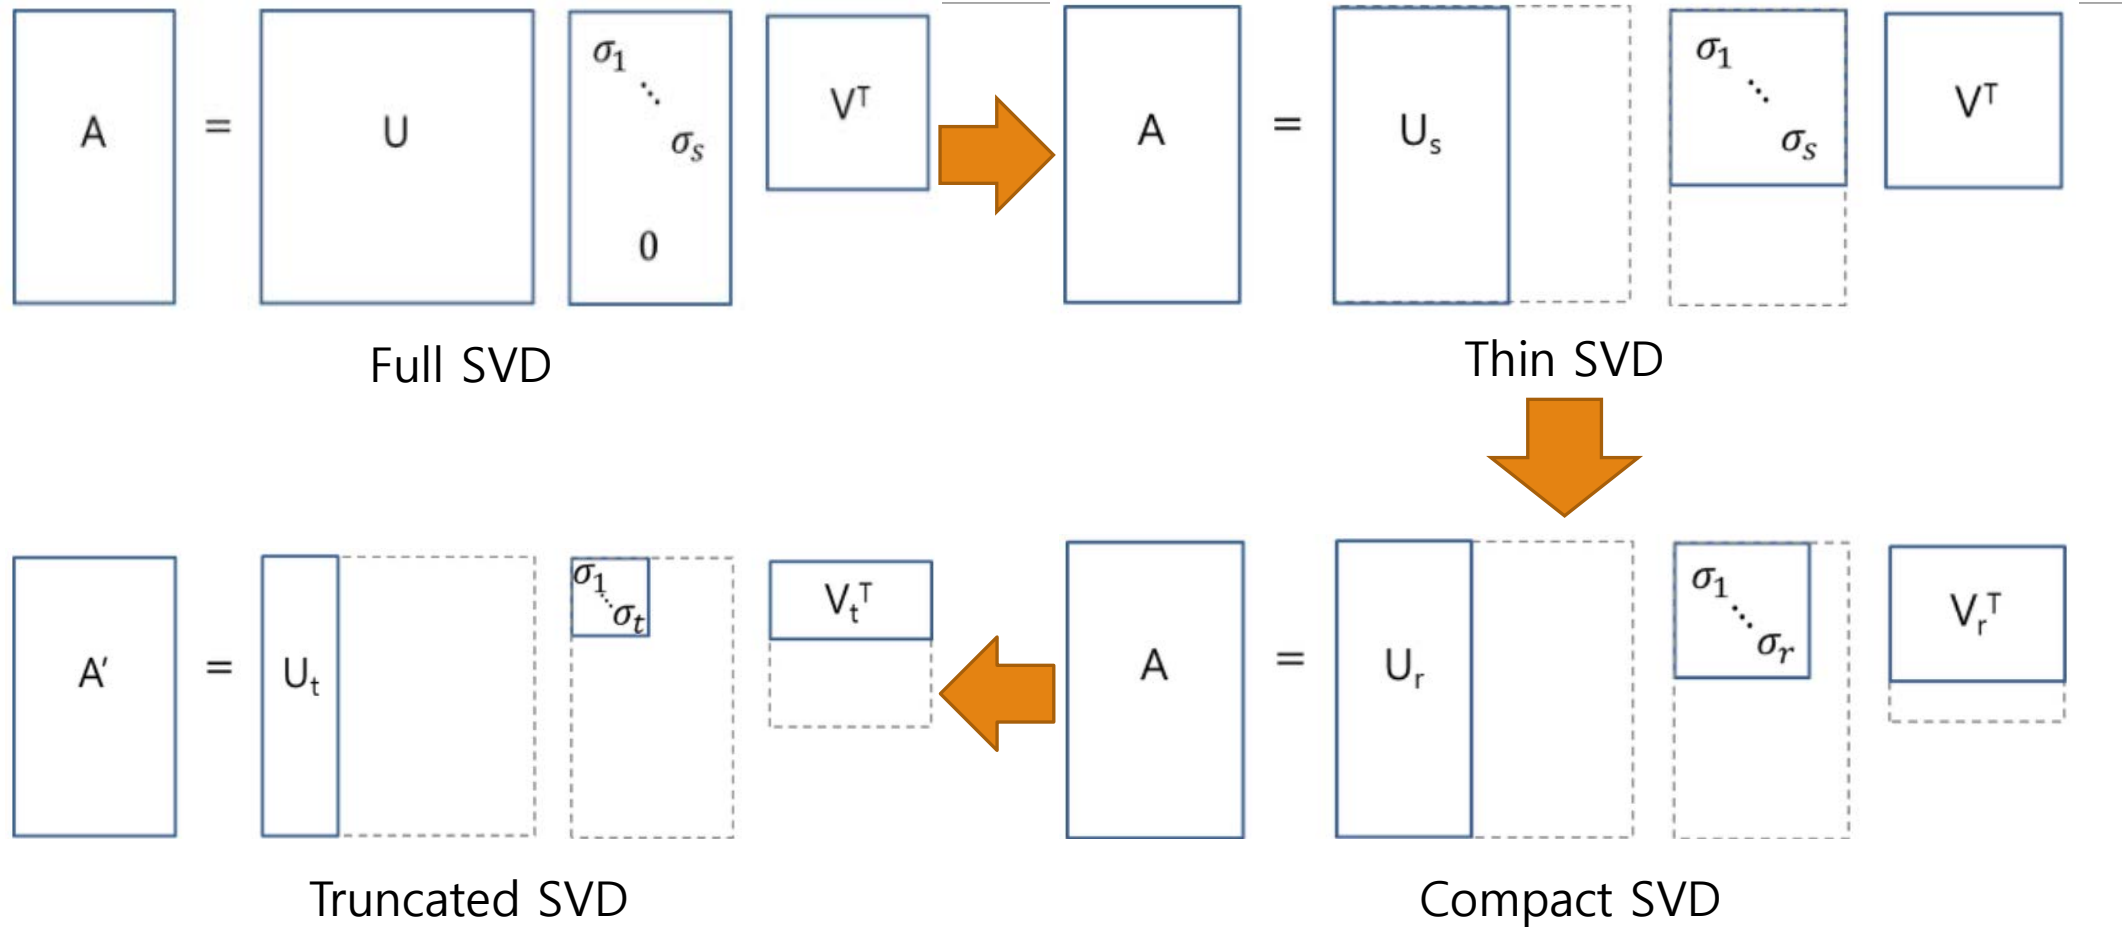


# Matrix rank constraint

---



# Truncated svd



# Truncated SVD

---

Original image: rank 402




Truncated svd: rank 80



# Canonical Polyadic decomposition

Weight for Convolutional layer: 4-tensor



# of parameters:  $T_1 T_2 T_3 T_4 \rightarrow R(T_1 + T_2 + T_3 + T_4)$



# Low rank approximation

Model	TOP-5 Accuracy	Speed-up	Compression Rate
AlexNet	80.03%	1.	1.
BN Low-rank	80.56%	1.09	4.94
CP Low-rank	79.66%	1.82	5.
VGG-16	90.60%	1.	1.
BN Low-rank	90.47%	1.53	2.72
CP Low-rank	90.31%	2.05	2.75
GoogleNet	92.21%	1.	1.
BN Low-rank	91.88%	1.08	2.79
CP Low-rank	91.79%	1.20	2.84

